

Document ingestion converts files into searchable text. For PDFs, page-level extraction preserves source and page numbers so later answers can cite evidence. Cleaning should remove extraction artifacts without rewriting the original wording.

Chunking splits long pages into overlapping windows so a sentence that crosses a boundary is not lost. Local embeddings turn each chunk into a vector. FAISS stores those vectors for similarity search. Metadata must keep a one-to-one map from vector position to chunk text.